Improved Loop Execution Modeling in the Clang Static Analyzer

Péter Szécsi

Eötvös Loránd University
Budapest, Hungary
peterszecsi95@inf.elte.hu

Abstract

The LLVM Clang Static Analyzer is a source code analysis tool which aims to find bugs in C, C++, and Objective-C programs using symbolic execution, i.e. it simulates the possible execution paths of the code. Currently the simulation of the loops is somewhat naive (but efficient), unrolling the loops a predefined constant number of times. However, this approach can result in loss of coverage in various cases.

This study aims to introduce two alternative approaches which can extend the current one and can be applied simultaneously: (1) loops worth to fully unroll were determined with applied heuristics (2) using a widening mechanism to simulate an arbitrary number of step of iteration. These methods were evaluated on numerous open source projects, and were proved to increase coverage in most of the cases. This work also laid the infrastructure for future loop modeling improvements.

Keywords   static analysis, symbolic execution, loop modeling

1. Introduction

The Clang Static Analyzer finds bugs by performing a symbolic execution on the code. During symbolic execution, the program is being interpreted, on a function-by-function basis, without any knowledge about the runtime environment. It builds up and traverses an inner model of the execution paths, called ExplodedGraph, for every analyzed function.

An important technical note that the building of the `ExplodedGraph` is based on the `Control Flow Graph` (CFG) of the functions. The `CFG` represents a source-level, intra-procedural control-flow of a statement. This statement can potentially be an entire function body, or just a single expression. The `CFG` consists of `CFGBlocks` which are simply containers of statements. The `CFGBlocks` essentially represents the basic blocks of the code but can contain some extra custom information. So for the sake of easier understanding TODO amikor azt mondjuk, hogy basic block, akkor CFGBlock-ra gondolunk, de ez segit elkepzelni/ jobban kifejezni/szemleletesebbe teszi a peldakat, ha basic blockot kepzelunk magunk ele.

A node of this graph (called ExplodedNode) contains a ProgramPoint (which determines the location) and a State (which contains the known information at that point). Its paths from the root to the leaves are modeling the different execution paths of the analyzed function. Whenever the execution encounters a branch, a corresponding branch will be created in the `ExplodedGraph` during the simulated interpretation. Hence, branches lead to an exponential number of ExplodedNodes. This combinatorial explosion is handled in the Static Analyzer by stopping the analysis when given conditions are fulfilled. Ceasing the analyzation process may cause loss of potential true positive results, but it is indispensable for maintaining a reasonable resource consumption regarding the memory and CPU usage. These conditions are modeled by the concept of budget. The budget is a collection of limitations on the shape of the `ExplodedGraph`. These limitations include:

1. The maximum number of traversed nodes in the `ExplodedGraph`. If this number is reached then the analysis of the simulated function stops.

2. The size of the simulated call stack. When a function call is reached then the analysis continues in its body as if it was inlined to the place of call (interprocedural). There are several heuristics that may control the behavior of inlining process. For example the too large functions are not inlined at all, and the really short functions are not counted in the size of call stack.

3. The number of times a function is inlined. The idea behind this constraint is that the more a function is analyzed, the less likely it is that a bug will appear in it. If this number is reached then that function will not be inlined again in this `ExplodedGraph`.

4. The number of times a basic block is processed during the analysis. This constraint limits the number of loop itera-

tions. When this threshold is reached the currently analyzed execution path will be aborted. The budget expression can be used in two ways. Sometimes it means the collection of the limitations above, sometimes it refers to one of these limitations. This will always be distinguishable from the context.

2. Motivation

Currently, the analyzer handles loops quite simply. More precisely, it unrolls them 4 times by default and then cuts the analysis of the path where the loop would have been unrolled more than 4 times. This behavior is reached by the above presented basic block visiting budget.

Loss in code coverage is one of the problems with this approach to loop modeling. Specifically, in cases where the loop is statically known to make more than 4 steps, the analyzer do not analyze the code following the loop. Thus, the naive loop handling (described above) could lead to entirely unchecked code. Listing 1 shows a small example for that.

```
1  void foo() {
2  int arr[6];
3  for (int i = 0; i < 6; i++){
4  arr[i] = i;
5  }
6  /* rest of the function */
7  }
```

Listing 1. Since the loop condition is known at every iteration, the analyzer will not check the 'rest of the function' part in the current state.

According to the budget rule concerning the basic block visit number, the analysis of the loop stops in the fourth iteration even if the loop condition is simple enough to see that unrolling the whole loop would not be too much extra work relatively. Running out of the budget implies (in this case) that the rest of the function body remains unanalyzed, which may lead to not finding potential bugs. Another problem can be seen on Listing 2:

```
1  int num();
2  void foo()
3  {
4  int n = 0;
5  for (int i = 0; i < num(); ++i) {
6  ++n;
7  }
8  /* rest of the function, n < 4 */
9  }
```

Listing 2. The loop condition is unknown but the analyzer will not generate a simulation path where n ≥ 4 (which can result coverage loss).

This code fragment results in an analysis which keeps track of the values of n and i variables (this information is stored in the State). In every iteration of the loop the values are updated accordingly. Note that updating the State means that a new node is inserted to the ExplodedGraph with the new values. Since the body of the num() function is unknown, the analyzer can not find out its return value. Thus it is considered as unknown. This circumstance makes the graph split into two branches. The first one belongs to the symbolic execution of the loop body assuming that the loop condition is true. The other one simulates the case where the condition is false and the execution continues after the loop. This process is done for every loop iteration, however, the 4th time assuming the condition is true, the path will be cut short according to the budget rule. Although the analyzer generates paths to simulate the code after the loop in the above described case, yet the value of variable n will be always less than 4 on these paths and the rest of the function will only be checked assuming this constraint. This can result in coverage loss as well, since the analyzer will ignore the paths where n is more than 4.

3. Proposed Solution

In this section two solutions are presented to resolve the previously mentioned limitations on symbolic execution of loops in the Clang Static Analyzer. It is important to note that these enhancements are incremental in the sense that on examples which are too complex to handle at the moment, we fall back to the original method. For the sake of simplicity in the following examples a "division by zero" will illustrate the bug we intend to find.

3.1 Loop Unrolling Heuristics

Loop unrolling means we have identified heuristics and patterns (such as loops with small number of branches and small known static bound) in order to find specific loops which are worth to be completely unrolled. This idea is inspired by the following example:

```
1  void foo() {
2  for (int i = 0; i < 6; i++){
3  /* simple loop which does not
4  change 'i' or split the state */
5  }
6  int k = 0;
7  int l = 2/k; // Division by zero
8  }
```

Listing 3. Complete unrolling of the loop makes it possible to find the division by zero error.

In the current solution a loop has to fulfill the following conditions in order to be unrolled:

1. The loop condition should arithmetically compare a variable – which is known at the beginning of the loop – to a literal (like: i < 6 or 6 ≥ i)

2. The loop should change the loop variable only once in its body and the difference needs to be constant. (This way the maximum number of steps can be estimated.)

3. The estimated number of steps should be less than 128. (Simulating loops which takes thousands of steps because they could single handedly exhaust the budget.)

4. The loop must not generate new branches or use `goto` statements.

By using this method, the bug on the Listing 3 example can be found successfully.

3.2 Loop Widening

The final aim of widening is quite the same as the unrolling, to increase the coverage of the analysis. However, it reaches it in a very different way. During widening the analyzer simulates the execution of an arbitrary number of iterations. The analyzer already has a widening algorithm, which reaches this behavior by discarding all of the known information before the last step of the loop. So, the analyzer creates the paths for the first 3 steps and simulate them as usual, but the widening (means the invalidating) happens before the 4th step in order to not lose the first precise simulation branches. This way the coverage will be increased, however, this method is disabled by default, since easily can result in too much false positives. Consider the example on Listing 4.

```
1  int num();
2  void foo() {
3  bool b = true;
4  for (int i = 0; i < num(); ++i) {
5  /*does not changes 'b'*/
6  }
7  int n = 0;
8  if (b)
9  n++;
10 n = 1/n; // False positive:
11 // Division by zero
12 }
13 }
```

Listing 4. Invalidating every known information (even the ones which are not modified by the loop) can easily result false positives.

In this case the analyzer will create and check that unfeasible path where the variable b is false, so n is not incremented and lead into a division by zero error. Since this execution path would never be performed while running the analyzed program, it is considered a false positive. My aim was to give a more precise approach for widening.

The principles are that we try to continue the analysis after the block visiting budget is exhausted but invalidate the information only on the variables which are possibly modified by the loop. For this I developed a solution which checks every possible way in which a variable can be modified in the loop. Then it evaluates these cases and if it encounters a modified variable which cannot be handled by the invalidation process (e.g.: a pointer variable) then the loop will not be widened and we return to the conservative method. This mechanism ensures that we do not create nodes containing

invalid states. This approach helps us to cover cases and find bugs like the one illustrated on Listing 5, and still not report false positives which are represented on Listing 4.

```
1  int num();
2  void foo() {
3  int n = 0;
4  for (int i = 0; i < num(); ++i) {
5  ++n;
6  }
7  if (n > 4) {
8  int k = 0;
9  k = 1/k;  // Division by zero error
10 }
11 }
```

Listing 5. Invalidating the information on only the possible changed variables can result higher coverage (while limiting the number of the found false positives).

We find the bug since we invalidate the known informations on variable n (and i as well). This cause the analyzer to create a branch where it checks the body of the if statement and finds the bug. However, this solution has its own limitations when dealing with nested loops. Consider the case on Listing 6.

```
1  int num();
2  void foo() {
3  int n = 0;
4  for (int i = 0; i < num(); ++i) {
5  ++n;
6  for (int j = 0; j < 4; ++j) {
7  /*body that does not change n*/
8  }
9  }
10 /*rest of the function, n <= 1 */
11 }
```

Listing 6. The naive widening method does not handle well the nested loops. In this example the outer loop will not be widened.

In this scenario, when the analyzer first step into the outer loop (so, assumes that i < num() is true) and encounter the inner loop, then it consumes its (own) block visiting budget. (This implies that it will be widened, although in this case it means that only the inner loop counter (j) information is discarded.) After that we move on to the next iteration, and assume that we are on the path where the outer loop condition is true again. Because we already exhausted the budget in the previous iteration, the next visit of the first basic block of the inner loop (the condition) means that this path will be completely cut off and not analyzed. This results that the outer loop will not reach the step number where it would been widened. Furthermore, the outer loop will not even reach the 3rd step and the 2nd is stopped at in its body as well (as described above). This causes the problem, that even though we use the loop widening method, we will analyze the rest of the function with the assumption n <= 1.

In order to deal with the above described nested loop problem, I have implemented a replay mechanism. This means that whenever we encounter an inner loop which already consumed its budget, we replay the analysis process of the current step of the outer loop but performing a widening first. This ensures the creation of a path which assumes that the condition is false and simulates the execution after the loop while the possibly changed information are discarded. This way the analyzer will not exclude some feasible path because of the simple loop handling which solves the problem.

Another note to the widening process that it makes sense to analyze the branch where the condition is true with the widened State as well. The example on Listing 7 shows a case where this is useful.

```
1  int num();
2  void foo() {
3    int n = 0;
4    int i;
5    for (i = 0; i < num(); ++i) {
6      if (i == 7) {
7        break;
8      }
9      for (int j = 0; j < 4; ++j) {/* */}
10   }
11   int n = 1 / (7 - k);
12          // ^ Possible division by zero
13 }
```

Listing 7. The replay mechanism successfully helps us to find the possible error the outer loop.

This way the analyzer will produce a path where the value of `i` is known to be 7, so it will be able find the possible division by zero error.

3.3 Infrastructure improvements

The discussed approaches heavily rely on informations that we are able to determine the following things:

1. Decide on any `ExplodedNode` if it simulates a body of a loop or not.
2. The point of entering and exiting a loop on the simulated path.

However, these information were not provided by the analyzer earlier.

4. Evaluation

The effect of the described loop modeling approaches were measured on various C/C++ open source project. These are the following:

Project	LoC	Language
TinyXml	20k	C++
Curl	21k	C
Redis	40k	C
Xerces	228k	C++
Vim	540k	C
OpenSSL	550k	C
PostgreSQL	950k	C
FFmpeg	1080k	C

4.1 Coverage and the number of explored paths

Project	Normal	Unroll	Widen	U + W
Curl	92.2	92.3	92.7	92.7
Redis	68.5	69.1	68.5	71.3
Xerces	92.2	92.3	92.7	92.7
Vim	1	2	3	4
OpenSSL	97.4	2	3	98.0
PostgreSQL	1	2	3	4
FFmpeg	1	2	3	4
Project	Normal	Unroll	Widen	U + W
Curl	92.2	92.3	92.7	92.7
Redis	68.5	69.1	68.5	71.3
Xerces	92.2	92.3	92.7	92.7
Vim	1	2	3	4
OpenSSL	97.4	2	3	98.0
PostgreSQL	1	2	3	4
FFmpeg	1	2	3	4

4.2 Analysis time

Project	Normal	Unroll	Widen	U + W
Curl	92.2	92.3	92.7	92.7
Redis	68.5	69.1	68.5	71.3
Xerces	92.2	92.3	92.7	92.7
Vim	1	2	3	4
OpenSSL	97.4	2	3	98.0
PostgreSQL	1	2	3	4
FFmpeg	1	2	3	4

4.3 Found bugs

Project	Normal	Unroll	Widen	U + W
Curl	92.2	92.3	92.7	92.7
Redis	68.5	69.1	68.5	71.3
Xerces	92.2	92.3	92.7	92.7
Vim	1	2	3	4
OpenSSL	97.4	2	3	98.0
PostgreSQL	1	2	3	4
FFmpeg	1	2	3	4

5. Conclusion

We introduced two alternative approaches for simulating loops during symbolic execution. These were implemented and subsequently evaluated on various open source projects, with a clear improvement of code coverage in general. With

the help of the new methods we are able to explore previously skipped, feasible execution paths, especially when they were used combined.

The required changes done to the underlying infrastructure should also ease implementing future enhancements. In particular, the information tracked by the analysis about the location contexts has been expanded with additional fields. While code coverage was measured to have increased by an average of *x*, there also was a noticeable performance penalty. In some cases, we observed a runtime increase of *y*, with an average of ... . The number of simulated paths also increased proportionally with the time taken, suggesting this time was well spent.

6. Future work

The heuristic patterns for completely unrolled loops could be extended to involve loops which bound is a known variable which is not changed in the body. Furthermore, even more general rules would be beneficial: consider loops where the value variables are known at the beginning and they are affected by a known constant change by every iteration. These improvements were not implemented yet due to some technical and framework limitations.

During the widening process we invalidate the possible changed informations. However, in case of a change made on a pointer could mean that we need to invalidate all the variables due to the lack of advanced pointer analysis. So, introducing pointer analysis algorithms to the analyzer could help to reach a more precise invalidation process.

The infrastructural improvements enable the analyzer to provide entry points for the bug finding modules (checkers) on loop entrances/exits and identify the currently simulated loop for every `ExplodedNode`. On top of these entry points new checkers can be implemented.

Acknowledgments

References

P. Q. Smith, and X. Y. Jones. ...reference text...